

Elasticsearch - Week 6

Assignment Report

Siddharth Kumar

September 15, 2025

Contents

1 Task 1: Elasticsearch 3-Node Cluster Setup	2
1.1 Cluster Setup and Configuration	2
1.2 Sub Tasks Execution	4
1.2.1 Create an Index	4
1.2.2 Insert Data into Index	5
1.2.3 Read Data and Export to JSON	7
1.2.4 Capture Cluster Health (Before Node Shutdown)	9
1.2.5 Stop Two Nodes and Verify Data Availability	10
1.2.6 Capture Cluster Health (After Node Shutdown)	12
1.2.7 Restart Nodes and Capture Cluster Health	13
2 Task 2: Monitoring and Alerting	14
2.1 Alerting Example	14

Chapter 1

Task 1: Elasticsearch 3-Node Cluster Setup

1.1 Cluster Setup and Configuration

```
1 #node1 ,2,3
2 wget https://artifacts.elastic.co/downloads/elasticsearch/
3   elasticsearch-9.1.3-linux-x86_64.tar.gz
4 wget https://artifacts.elastic.co/downloads/elasticsearch/
5   elasticsearch-9.1.3-linux-x86_64.tar.gz.sha512
6 shasum -a 512 -c elasticsearch-9.1.3-linux-x86_64.tar.gz.
7   sha512
8 tar -xzf elasticsearch-9.1.3-linux-x86_64.tar.gz
9 cd elasticsearch-9.1.3/
10
11 #node1
12 #/home/sre/elasticsearch-9.1.3/config/elasticsearch.yml
13 cluster.name: ppe_elasticsearch
14 node.name: node1
15 node.roles: [ master, data ]
16 network.host: 0.0.0.0
17 transport.host: 0.0.0.0
18 http.port: 9200
19 cluster.initial_master_nodes: ["node1"]
20
21 #shell
22 ./bin/elasticsearch -d -p pid
23 ./bin/elasticsearch-create-enrollment-token -s node
24 $ES_HOME/bin/elasticsearch-reset-password -u elastic
```

Listing 1.1: Elasticsearch node configuration

Explanation

- Line 1: Indicates that these steps are to be repeated on node1, node2, and node3.
- Lines 2–3: Download the Elasticsearch 9.1.3 package and its checksum file from Elastic’s repository.
- Line 5: Extract the Elasticsearch archive.
- Line 6: Navigate into the extracted installation directory.
- Lines 9–16: Edit `elasticsearch.yml`: set cluster name, unique node name, roles (master+data), bind network/transport to all interfaces, configure HTTP port, and define initial master node(s).
- Line 19: Start Elasticsearch as a background service and save its PID.
- Line 20: Generate an enrollment token so additional nodes can securely join the cluster.
- Line 21: Reset the built-in `elastic` user password for authentication.

```
1 #node2-3
2 #/home/sre/elasticsearch-9.1.3/config/elasticsearch.yml
3 cluster.name: ppe_elasticsearch
4 node.name: node1
5 node.roles: [ master, data ]
6
7 #shell
8 ./bin/elasticsearch --enrollment-token <enrollment-token>
```

Listing 1.2: Elasticsearch node2-3 configuration

```
1 curl -u elastic:$ELASTIC_PASSWORD https://10.57.40.168:9200/_cat/nodes?v -k
```

Listing 1.3: API call to check cluster nodes

```
[sre@stg-hdpsiddharth101:~/elasticsearch-9.1.3$ curl -u elastic:$ELASTIC_PASSWORD https://10.57.40.168:9200/_cat/nodes?v -k
ip          heap.percent ram.percent cpu load_1m load_5m load_15m node.role master name
10.57.40.168      9        65   0   0.15   0.05   0.01 dm     *    node1
10.57.40.169      2        65   0   0.03   0.07   0.02 dm     -    node3
10.57.40.170     14        65   0   0.07   0.04   0.00 dm     -    node2
...
```

Figure 1.1: Cluster nodes

1.2 Sub Tasks Execution

1.2.1 Create an Index

```
1 curl -u elastic:<password> -X PUT "https://10.57.40.168:9200/ppe_index?pretty" -k
```

Listing 1.4: Create index ppe_index

```
[sre@stg-hdpsiddharth101:~/elasticsearch-9.1.3$ curl -u elastic:$ELASTIC_PASSWORD -X PUT "https://10.57.40.168:9200/ppe_index?pretty" -k
{
  "acknowledged" : true,
  "shards_acknowledged" : true,
  "index" : "ppe_index"
}
```

Figure 1.2: Index creation successfull

```
[sre@stg-hdpsiddharth101:~/elasticsearch-9.1.3$ curl -u elastic:$ELASTIC_PASSWORD -k "https://10.57.40.168:9200/_cat/shards?v"
index           shard prirep state  docs store dataset ip      node
.security-7          0   p    STARTED  35 84.9kb 84.9kb 10.57.40.170 node2
.security-7          0   r    STARTED  35 84.9kb 84.9kb 10.57.40.169 node3
.ds-ilm-history-7-2025.09.13-000001 0   p    STARTED  3  9.8kb  9.8kb 10.57.40.168 node1
.ds-ilm-history-7-2025.09.13-000001 0   r    STARTED  3  9.8kb  9.8kb 10.57.40.169 node3
ppe_index            0   p    STARTED  0   249b   249b 10.57.40.168 node1
ppe_index            0   r    STARTED  0   249b   249b 10.57.40.168 node2
.ds-logs-elasticsearch.deprecation-default-2025.09.13-000001 0   r    STARTED  1 10.3kb 10.3kb 10.57.40.170 node2
.ds-logs-elasticsearch.deprecation-default-2025.09.13-000001 0   p    STARTED  1 10.3kb 10.3kb 10.57.40.168 node1
```

Figure 1.3: Default shards allocation

1.2.2 Insert Data into Index

```
1 #!/bin/bash
2
3 ES_HOST="https://10.57.40.168:9200"
4 ES_USER="elastic"
5 ES_PASS="4+fgoJIF16HbF4WxAFlR"
6 INDEX="ppe_index"
7
8 echo "Inserting 30 documents into index: $INDEX"
9
10 for i in {1..30}; do
11     curl -s -u $ES_USER:$ES_PASS -k -X POST "$ES_HOST/$INDEX/
12         _doc?pretty" \
13             -H 'Content-Type: application/json' \
14             -d "{\"doc_id\": $i,
15                 \"user\": { \"id\": \"siddharth\" },
16                 \"message\": \"Document number $i\""
17             }" > /dev/null
18     echo "Inserted doc_id=$i"
19 done
20
21 echo "Finished inserting documents."
```

Listing 1.5: Insert documents script

Explanation

- Line 1: Declares the script as a Bash script using the shebang.
- Lines 3–6: Define variables for Elasticsearch host URL, username, password, and the index name (`ppe_index`).
- Line 8: Prints a message indicating that documents are about to be inserted.
- Line 10: Starts a loop from 1 to 30 to generate 30 documents.
- Line 11: Uses `curl` with authentication to send a POST request to Elasticsearch for each document. The `-k` flag allows insecure SSL connections.
- Line 12: Sets the request header to `Content-Type: application/json`.
- Lines 13–15: Define the JSON body of the document: each document has a unique `doc_id`, a user object with id `siddharth`, and a message field with the document number.
- Line 16: Redirects `curl` output to `/dev/null` to avoid cluttering the console.
- Line 17: Prints a confirmation message for each inserted document.
- Line 18: Ends the loop after inserting 30 documents.
- Line 20: Prints a final message once all documents have been inserted.

```
[sre@stg-hdpsiddharth101:~/elasticsearch-9.1]$ ./insert_docs.sh
✖ Inserting 30 documents into index: ppe_index
Inserted doc_id=1
Inserted doc_id=2
Inserted doc_id=3
Inserted doc_id=4
Inserted doc_id=5
Inserted doc_id=6
Inserted doc_id=7
Inserted doc_id=8
Inserted doc_id=9
Inserted doc_id=10
Inserted doc_id=11
Inserted doc_id=12
Inserted doc_id=13
Inserted doc_id=14
Inserted doc_id=15
Inserted doc_id=16
Inserted doc_id=17
Inserted doc_id=18
Inserted doc_id=19
Inserted doc_id=20
Inserted doc_id=21
Inserted doc_id=22
Inserted doc_id=23
Inserted doc_id=24
Inserted doc_id=25
Inserted doc_id=26
Inserted doc_id=27
Inserted doc_id=28
Inserted doc_id=29
Inserted doc_id=30
✓ Finished inserting documents.
[sre@stg-hdpsiddharth101:~/elasticsearch-9.1.3$ curl -s -u elastic:$ELASTIC_PASSWORD -k https://10.57.40.168:9200/ppe_index/_count?pretty
{
  "count" : 30,
  "_shards" : {
    "total" : 1,
    "successful" : 1,
    "skipped" : 0,
    "failed" : 0
  }
}
```

Figure 1.4: Script inserting docs

1.2.3 Read Data and Export to JSON

```
1 curl -u elastic:<password> -X GET "https://10.57.40.168:9200/
  2   ppe_index/_search?pretty&size=1000" -k \
    -H 'Content-Type: application/json' > read_data.json
```

Listing 1.6: Read data and write to JSON

```

    "took" : 178,
    "timed_out" : false,
    "_shards" : {
        "total" : 1,
        "successful" : 1,
        "skipped" : 0,
        "failed" : 0
    },
    "hits" : {
        "total" : {
            "value" : 30,
            "relation" : "eq"
        },
        "max_score" : 1.0,
        "hits" : [
            {
                "_index" : "ppe_index",
                "_id" : "6tzRR5kBFiGH02qahLh1",
                "_score" : 1.0,
                "_source" : {
                    "doc_id" : 1,
                    "user" : {
                        "id" : "siddharth"
                    },
                    "message" : "Document number 1"
                }
            },
            {
                "_index" : "ppe_index",
                "_id" : "69zRR5kBFiGH02qahbh_",
                "_score" : 1.0,
                "_source" : {
                    "doc_id" : 2,
                    "user" : {
                        "id" : "siddharth"
                    },
                    "message" : "Document number 2"
                }
            },
            {
                "_index" : "ppe_index",
                "_id" : "7NzRR5kBFiGH02qahbjF",
                "_score" : 1.0,
                "_source" : {
                    "doc_id" : 3,
                    "user" : {
                        "id" : "siddharth"
                    },
                    "message" : "Document number 3"
                }
            },
            {
                "_index" : "ppe_index",
                "_id" : "7dzRR5kBFiGH02qahrgI",
                "_score" : 1.0,
                "_source" : {
                    "doc_id" : 4,
                    "user" : {
                        "id" : "siddharth"
                    },
                    "message" : "Document number 4"
                }
            },
            {
                "_index" : "ppe_index",
                "_id" : "7tzRR5kBFiGH02qahrhH",
                "_score" : 1.0,
                "_source" : {
                    "doc_id" : 5,
                    "user" : {
                        "id" : "siddharth"
                    },
                    "message" : "Document number 5"
                }
            }
        ]
    }
}

"read_data.json" 379L, 8588C

```

Figure 1.5: read_data.json

1.2.4 Capture Cluster Health (Before Node Shutdown)

```
1 #!/bin/bash
2
3 ES_HOST="https://10.57.40.168:9200"
4 ES_USER="elastic"
5 ES_PASS="4+fgoJIF16HbF4WxAFlR"
6 OUTPUT_FILE="cluster_health_before_step5.json"
7
8 # Get current timestamp
9 TIMESTAMP=$(date +"%Y-%m-%d %H:%M:%S")
10
11 # Capture cluster health and add timestamp
12 curl -s -u $ES_USER:$ES_PASS -k "$ES_HOST/_cluster/health/
    ppe_index?pretty" | \
13 jq --arg ts "$TIMESTAMP" '{timestamp: $ts} .' >
    $OUTPUT_FILE
14
15 echo "Cluster health saved to $OUTPUT\_FILE with timestamp:
    $TIMESTAMP"
```

Listing 1.7: Script to capture cluster health

Explanation

- Line 1: Declares the script as a Bash script with the shebang.
- Lines 3–5: Define variables for Elasticsearch host URL, username, and password.
- Line 6: Sets the output filename as `cluster_health_before_step5.json`.
- Line 9: Captures the current timestamp in YYYY-MM-DD HH:MM:SS format using the `date` command.
- Lines 12–13: Sends a request to Elasticsearch's `_cluster/health` API for the index `ppe_index`, using `curl` with authentication.
 - The `-k` flag allows insecure SSL connections.
 - The response is piped to `jq`, which injects the captured timestamp as a new JSON field.
 - The result is written to the output file defined earlier.
- Line 15: Prints a confirmation message showing the filename and timestamp of the saved cluster health data.

```
{  
  "timestamp": "2025-09-14 17:06:24",  
  "cluster_name": "ppe_elasticsearch",  
  "status": "green",  
  "timed_out": false,  
  "number_of_nodes": 3,  
  "number_of_data_nodes": 3,  
  "active_primary_shards": 1,  
  "active_shards": 2,  
  "relocating_shards": 0,  
  "initializing_shards": 0,  
  "unassigned_shards": 0,  
  "unassigned_primary_shards": 0,  
  "delayed_unassigned_shards": 0,  
  "number_of_pending_tasks": 0,  
  "number_of_in_flight_fetch": 0,  
  "task_max_waiting_in_queue_millis": 0,  
  "active_shards_percent_as_number": 100  
}  
~  
~
```

Figure 1.6: Cluster Health Before Node Shutdown

1.2.5 Stop Two Nodes and Verify Data Availability

```
1 #node2 -3  
2 pkill -f elasticsearch
```

Listing 1.8: Stop nodes and verify data

Explanation

After stopping two nodes, quorum fails, no master available. But read is working fine from replica on node1.

```

    "took" : 53,
    "timed_out" : false,
    "_shards" : {
        "total" : 1,
        "successful" : 1,
        "skipped" : 0,
        "failed" : 0
    },
    "hits" : [
        {
            "total" : {
                "value" : 30,
                "relation" : "eq"
            },
            "max_score" : 1.0,
            "hits" : [
                {
                    "_index" : "ppe_index",
                    "_id" : "6tzRR5kBFIgh02qahLh1",
                    "_score" : 1.0,
                    "_source" : {
                        "doc_id" : 1,
                        "user" : {
                            "id" : "siddharth"
                        },
                        "message" : "Document number 1"
                    }
                },
                {
                    "_index" : "ppe_index",
                    "_id" : "69zRR5kBFIgh02qahbh_",
                    "_score" : 1.0,
                    "_source" : {
                        "doc_id" : 2,
                        "user" : {
                            "id" : "siddharth"
                        },
                        "message" : "Document number 2"
                    }
                },
                {
                    "_index" : "ppe_index",
                    "_id" : "7NzRR5kBFIgh02qahbjF",
                    "_score" : 1.0,
                    "_source" : {
                        "doc_id" : 3,
                        "user" : {
                            "id" : "siddharth"
                        },
                        "message" : "Document number 3"
                    }
                },
                {
                    "_index" : "ppe_index",
                    "_id" : "7dzRR5kBFIgh02qahrgI",
                    "_score" : 1.0,
                    "_source" : {
                        "doc_id" : 4,
                        "user" : {
                            "id" : "siddharth"
                        },
                        "message" : "Document number 4"
                    }
                },
                {
                    "_index" : "ppe_index",
                    "_id" : "7tzRR5kBFIgh02qahrhH",
                    "_score" : 1.0,
                    "_source" : {
                        "doc_id" : 5,
                        "user" : {
                            "id" : "siddharth"
                        },
                        "message" : "Document number 5"
                    }
                }
            ]
        }
    ],
    "read_data_after_step5.json" 379L, 8587C

```

Figure 1.7: read_data_after_step5.json

1.2.6 Capture Cluster Health (After Node Shutdown)

Explanation

_cluster/health API not working because no master is available for that.Hence we are not able to get status of cluster.

```
{  
    "timestamp": "2025-09-14 17:22:49",  
    "error": {  
        "root_cause": [  
            {  
                "type": "master_not_discovered_exception",  
                "reason": null  
            }  
        ],  
        "type": "master_not_discovered_exception",  
        "reason": null  
    },  
    "status": 503  
}  
~  
~  
~
```

Figure 1.8: Cluster Health after Node Shutdown

1.2.7 Restart Nodes and Capture Cluster Health

```
1 #on node2 and node 3
2 $ES_HOME/bin/elasticsearch
```

Listing 1.9: Restart nodes

Explanation

After restarting, qorum is formed and master is elected. replica on node 1 is upgraded to primary. and it's repilca is made on node3.

```
[{"timestamp": "2025-09-14 17:31:43",
"cluster_name": "ppe_elasticsearch",
"status": "green",
"timed_out": false,
"number_of_nodes": 3,
"number_of_data_nodes": 3,
"active_primary_shards": 1,
"active_shards": 2,
"relocating_shards": 0,
"initializing_shards": 0,
"unassigned_shards": 0,
"unassigned_primary_shards": 0,
"delayed_unassigned_shards": 0,
"number_of_pending_tasks": 0,
"number_of_in_flight_fetch": 0,
"task_max_waiting_in_queue_millis": 0,
"active_shards_percent_as_number": 100
}
~
~
~
```

Figure 1.9: Cluster Health after restarting Nodes

```
[sre@stg-hdpsiddharth101:~/elasticsearch-9.1.3$ curl -s -u elastic:$ELASTIC_PASSWORD -k https://localhost:9200/_cat/shards/?v -k
index shard prirep state docs store dataset ip          node
.security-7      0   p   STARTED 35 84.9kb 84.9kb 10.57.40.170 node2
.security-7      0   r   STARTED 35 84.9kb 84.9kb 10.57.40.169 node3
.ds-ilm-history-7-2025.09.13-000001 0   p   STARTED 3 9.8kb 9.8kb 10.57.40.168 node1
.ds-ilm-history-7-2025.09.13-000001 0   r   STARTED 3 9.8kb 9.8kb 10.57.40.169 node3
.ppe_index        0   p   STARTED 30 13.8kb 13.8kb 10.57.40.168 node1
.ppe_index        0   r   STARTED 30 13.8kb 13.8kb 10.57.40.169 node3
.ds-logs-elasticsearch.deprecation-default-2025.09.13-000001 0   r   STARTED 1 10.3kb 10.3kb 10.57.40.170 node2
.ds-logs-elasticsearch.deprecation-default-2025.09.13-000001 0   p   STARTED 1 10.3kb 10.3kb 10.57.40.168 node1
```

Figure 1.10: Shards allocation after restarting Nodes

Chapter 2

Task 2: Monitoring and Alerting

2.1 Alerting Example

```
1 #!/bin/bash
2
3 ALERT_FILE="/home/sre/elasticsearch-9.1.3/es_alerts"
4 ES_URL="https://localhost:9200"
5 USER="elastic"
6 PASS="4+fgoJIF16HbF4WxAFlR"
7
8 while true; do
9     TIMESTAMP=$(date '+%Y-%m-%d %H:%M:%S')
10
11     RESPONSE=$(curl -s -k -u "$USER:$PASS" "$ES_URL/_cluster/
12         health" 2>/dev/null)
13     CURL_EXIT=$?
14
15     if [ $CURL_EXIT -ne 0 ] || [ -z "$RESPONSE" ]; then
16         echo "[${TIMESTAMP}] ALERT: Cannot reach cluster or no
17             master elected!" >> "$ALERT_FILE"
18     else
19         STATUS=$(echo "$RESPONSE" | jq -r '.status // "null",
20             ')
21         UNASSIGNED=$(echo "$RESPONSE" | jq -r '.unassigned_shards // 0')
22         NODES=$(echo "$RESPONSE" | jq -r '.number_of_nodes //
23             0')
24
25         # Convert "null" to safe integers
26         [ "$UNASSIGNED" = "null" ] && UNASSIGNED=0
27         [ "$NODES" = "null" ] && NODES=0
```

```

24
25     if [ "$STATUS" != "green" ]; then
26         echo "[${TIMESTAMP}] ALERT: Cluster status is
27             $STATUS" >> "$ALERT_FILE"
28     fi
29
30     if [ "$UNASSIGNED" -gt 0 ]; then
31         echo "[${TIMESTAMP}] ALERT: Unassigned shards:
32             $UNASSIGNED" >> "$ALERT_FILE"
33     fi
34
35     if [ "$NODES" -lt 3 ]; then
36         echo "[${TIMESTAMP}] ALERT: Number of nodes is less
37             than 3: $NODES" >> "$ALERT_FILE"
38     fi
39
done
sleep 30

```

Listing 2.1: Script for alert

Explanation

- Lines 3–6: Define variables for alert log file path, Elasticsearch URL, and authentication credentials.
- Line 8: Begins an infinite loop that continuously checks the cluster.
- Line 9: Captures the current timestamp in YYYY-MM-DD HH:MM:SS format.
- Line 11: Uses curl to query the cluster health endpoint, saving the response into RESPONSE. Errors are redirected to /dev/null.
- Line 12: Stores curl's exit code in CURL_EXIT for connectivity checks.
- Lines 14–15: If curl failed or the response is empty, logs an alert indicating the cluster is unreachable or no master is elected.
- Lines 17–19: Parse JSON response with jq to extract status, number of unassigned_shards, and number_of_nodes.
- Lines 22–23: Replace any null values with safe integers (0).
- Lines 25–27: If the cluster status is not green, log an alert with the current status.
- Lines 29–31: If there are unassigned shards, log an alert with their count.
- Lines 33–35: If the number of nodes is less than 3, log an alert showing the current node count.
- Line 38: Waits 30 seconds before the next check.
- Line 39: Ends the infinite loop, ensuring continuous monitoring.

```
[2025-09-14 22:11:19] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:11:49] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:12:20] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:12:50] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:13:20] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:13:50] ALERT: Cluster status is 503
[2025-09-14 22:14:51] ALERT: Cluster status is 503
[2025-09-14 22:15:51] ALERT: Cluster status is 503
[2025-09-14 22:16:51] ALERT: Cluster status is yellow
[2025-09-14 22:16:51] ALERT: Unassigned shards: 1
[2025-09-14 22:16:51] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:16:51] ALERT: Cluster status is yellow
[2025-09-14 22:17:21] ALERT: Unassigned shards: 1
[2025-09-14 22:17:21] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:17:52] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:18:22] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:18:52] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:19:23] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:19:53] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:20:23] ALERT: Number of nodes is less than 3: 2
[2025-09-14 22:20:53] ALERT: Number of nodes is less than 3: 2
~
~
```

Figure 2.1: Alert logs stored in es_alerts file