

Kafka Assignment

Siddharth Kumar

October 2, 2025

Contents

1	Task 1: Kafka Cluster Setup and Topic Operations	2
1.1	Topic Creation	2
1.2	Topic Description	3
1.3	Produce and Consume Data	4
1.4	Alter Partitions	5
1.5	Verifying data is produced to newer partitions	6
1.6	Stop Broker and Check Cluster Health	7
1.7	Again, produce and consume data for the topic test_topic	8
1.8	Describe the topic again and observe the changes	8
1.9	Under-Replicated Partitions	9
1.10	Again, start the Kaka broker and observe the health	10
1.11	Experimenting with test_topic_1 with a replication factor of 1	11
2	Task 2: Kafka Health Alert Setup	13
2.1	Bash Script for Alerts	13
2.2	Cron Job Setup	14

Chapter 1

Task 1: Kafka Cluster Setup and Topic Operations

1.1 Topic Creation

```
1 #path-> /usr/odp/3.2.2.0-1/kafka/bin  
2 ./kafka-topics.sh --create --bootstrap-server stg-xxx:6667 \  
3 --replication-factor 3 --partitions 10 --topic test_topic
```

Explanation

- **Line 2:** Executes the `./kafka-topics.sh` script with the `--create` action, connecting to the Kafka cluster via the specified `--bootstrap-server` at `stg-xxx:6667`.
- **Line 3:** Defines the topic properties: sets the `--replication-factor` to 3 for fault tolerance, divides the topic into 10 `--partitions` for parallelism, and names it `--topic test_topic`.

```
[root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --create --zookeeper stg-hdpsiddharth102:2181 --replication-factor 3 --partitions 10 --topic test_topic  
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore ('_') could collide. To avoid issues it is best to use either, but not both.  
Created topic test_topic.
```

Figure 1.1: Kafka topic creation output

1.2 Topic Description

```
1 #path-> /usr/odp/3.2.2.0-1/kafka/bin  
2 ./kafka-topics.sh --describe --bootstrap-server stg-xxx:6667 --  
topic test_topic
```

Explanation

- **Line 2:** Executes the `./kafka-topics.sh` script with the `--describe` action to fetch details about an existing topic. It connects to the cluster via the `--bootstrap-server` `stg-xxx:6667` and specifies the target topic as `--topic test_topic`.

```
^Croot@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdpsiddharth102:6667  
topic test_topic  
Topic: test_topic      TopicId: 9pFEK8R2SD20cXH1hQ9dgg PartitionCount: 10      ReplicationFactor: 3      Configs: compression.type=producer,  
insync.replicas=1,segment.bytes=1073741824,max.message.bytes=1000000,index.interval.bytes=4096,retention.bytes=-1,segment.index.bytes=1048576  
Topic: test_topic      Partition: 0      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1001,1003  
Topic: test_topic      Partition: 1      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002,1001  
Topic: test_topic      Partition: 2      Leader: 1001      Replicas: 1001,1003,1002      Isr: 1001,1003,1002  
Topic: test_topic      Partition: 3      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003,1001  
Topic: test_topic      Partition: 4      Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1001,1002  
Topic: test_topic      Partition: 5      Leader: 1001      Replicas: 1001,1002,1003      Isr: 1001,1002,1003  
Topic: test_topic      Partition: 6      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1001,1003  
Topic: test_topic      Partition: 7      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002,1001  
Topic: test_topic      Partition: 8      Leader: 1001      Replicas: 1001,1003,1002      Isr: 1001,1003,1002  
Topic: test_topic      Partition: 9      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003,1001
```

Figure 1.2: Kafka describe

1.3 Produce and Consume Data

```
1 # Produce data
2 #path-> /usr/odp/3.2.2.0-1/kafka/bin
3 ./kafka-console-producer.sh \
4 > --bootstrap-server stg-XXX:6667 \
5 > --topic test_topic \
6 > --property parse.key=true \
7 > --property key.separator=":" \
8
9
10 # Consume data
11 #path-> /usr/odp/3.2.2.0-1/kafka/bin
12 ./kafka-console-consumer.sh \
13 > --bootstrap-server stg-xxx:6667 \
14 > --topic test_topic \
15 > --from-beginning \
16 > --group test_consumer_group \
17 > --property print.key=true \
18 > --property print.partition=true \
19 > --property print.offset=true \
20 > --property print.timestamp=true
```

Explanation

Produce Data (Lines 3-7):

- **Line 3:** Executes the `./kafka-console-producer.sh` script, which starts an interactive command-line tool to send messages to a Kafka topic.
- **Line 4-5:** Connects to the cluster via `--bootstrap-server stg-XXX:6667` and specifies that messages should be sent to the `--topic test_topic`.
- **Line 6-7:** Configures the producer's properties. `parse.key=true` tells it to split input into a key and a value, and `key.separator=":"` sets the colon as the delimiter between the key and the value.

Consume Data (Lines 12-20):

- **Line 12:** Executes the `./kafka-console-consumer.sh` script, which starts a tool to read messages from a Kafka topic and print them to the console.
- **Line 13-14:** Connects to the same Kafka cluster and specifies it will read from `--topic test_topic`.
- **Line 15:** The `--from-beginning` flag instructs the consumer to read all existing messages in the topic, starting from the earliest offset.
- **Line 16:** Assigns this consumer to the `--group test_consumer_group`. Kafka uses consumer groups to track which messages have been read (offsets).
- **Line 17-20:** Sets properties to enrich the output. These flags instruct the consumer to print not just the message value, but also its metadata: the `key`, `partition`, `offset`, and `timestamp`.

```

root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-console-producer.sh \
> --bootstrap-server stg-hdpsiddharth102:6667 \
> --topic test_topic \
> --property parse.key=true \
> --property key.separator=":" \
>user1:hello world
>user2:hello sid
>user1:hhiiii
>user3:qwerty
>user2:poiuyt
>

```

Figure 1.3: Kafka producer

```

root@stg-hdpsiddharth104:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-console-consumer.sh \
> --bootstrap-server stg-hdpsiddharth102:6667 \
> --topic test_topic \
> --from-beginning \
> --group test_consumer_group \
> --property print.key=true \
> --property print.partition=true \
> --property print.offset=true \
> --property print.timestamp=true
CreateTime:1759214670283      Partition:4      Offset:0      user1    hello world
CreateTime:1759214691312      Partition:9      Offset:0      user2    hello sid
CreateTime:1759214704874      Partition:4      Offset:1      user1    hhiiii
CreateTime:1759214749138      Partition:5      Offset:0      user3    qwerty
CreateTime:1759214767365      Partition:9      Offset:1      user2    poiuyt

```

Figure 1.4: Kafka consumer

1.4 Alter Partitions

```

1 #path-> /usr/odp/3.2.2.0-1/kafka/bin
2 ./kafka-topics.sh --alter \
3 > --bootstrap-server stg-xxx:9092 \
4 > --topic test_topic --partitions 20

```

Explanation

- **Line 2:** Executes the `./kafka-topics.sh` script with the `--alter` action. This command is used to modify the properties of an existing topic, such as increasing its partition count.

```

root@stg-hdfsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdfsiddharth102:6667 -
ic test_topic
Topic: test_topic      TopicId: 9pFEK8R2SD20cXH1hQ9dgg PartitionCount: 20      ReplicationFactor: 3      Configs: compression.type=producer,
insync.replicas=1,segment.bytes=1073741824,max.message.bytes=1000000,index.interval.bytes=4096,retention.bytes=-1,segment.index.bytes=10485
    Topic: test_topic      Partition: 0      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1001,1003
    Topic: test_topic      Partition: 1      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002,1001
    Topic: test_topic      Partition: 2      Leader: 1001      Replicas: 1001,1003,1002      Isr: 1001,1003,1002
    Topic: test_topic      Partition: 3      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003,1001
    Topic: test_topic      Partition: 4      Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1001,1002
    Topic: test_topic      Partition: 5      Leader: 1001      Replicas: 1001,1002,1003      Isr: 1001,1002,1003
    Topic: test_topic      Partition: 6      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1001,1003
    Topic: test_topic      Partition: 7      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002,1001
    Topic: test_topic      Partition: 8      Leader: 1001      Replicas: 1001,1003,1002      Isr: 1001,1003,1002
    Topic: test_topic      Partition: 9      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003,1001
    Topic: test_topic      Partition: 10     Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002,1001
    Topic: test_topic      Partition: 11     Leader: 1001      Replicas: 1001,1003,1002      Isr: 1001,1003,1002
    Topic: test_topic      Partition: 12     Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003,1001
    Topic: test_topic      Partition: 13     Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1001,1002
    Topic: test_topic      Partition: 14     Leader: 1001      Replicas: 1001,1002,1003      Isr: 1001,1002,1003
    Topic: test_topic      Partition: 15     Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1001,1003
    Topic: test_topic      Partition: 16     Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002,1001
    Topic: test_topic      Partition: 17     Leader: 1001      Replicas: 1001,1003,1002      Isr: 1001,1003,1002
    Topic: test_topic      Partition: 18     Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003,1001
    Topic: test_topic      Partition: 19     Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1001,1002

```

Figure 1.5: Kafka describe

1.5 Verifying data is produced to newer partitions

```

1 #path-> /usr/odp/3.2.2.0-1/kafka/bin
2 ./kafka-consumer-groups.sh \
3 > --bootstrap-server stg-xxx:6667 \
4 > --describe \
5 > --group test_consumer_group

```

```

root@stg-hdfsiddharth104:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-consumer-groups.sh --bootstrap-server stg-hdfsiddharth102:6667 \
> --describe \
> --group test_consumer_group

Consumer group 'test_consumer_group' has no active members.

GROUP          TOPIC      PARTITION  CURRENT-OFFSET  LOG-END-OFFSET  LAG      CONSUMER-ID      HOST      CLIENT-ID
test_consumer_group test_topic  15          0              0              0      -          -          -
test_consumer_group test_topic  14          0              0              0      -          -          -
test_consumer_group test_topic  13          0              0              0      -          -          -
test_consumer_group test_topic  12          0              0              0      -          -          -
test_consumer_group test_topic  19          0              0              0      -          -          -
test_consumer_group test_topic  18          1              1              0      -          -          -
test_consumer_group test_topic  17          1              1              0      -          -          -
test_consumer_group test_topic  16          0              0              0      -          -          -
test_consumer_group test_topic  7           0              0              0      -          -          -
test_consumer_group test_topic  6           0              0              0      -          -          -
test_consumer_group test_topic  5           1              1              0      -          -          -
test_consumer_group test_topic  4           2              2              0      -          -          -
test_consumer_group test_topic  11          0              0              0      -          -          -
test_consumer_group test_topic  10          0              0              0      -          -          -
test_consumer_group test_topic  9           4              4              0      -          -          -
test_consumer_group test_topic  8           0              0              0      -          -          -
test_consumer_group test_topic  3           1              1              0      -          -          -
test_consumer_group test_topic  2           1              1              0      -          -          -
test_consumer_group test_topic  1           1              1              0      -          -          -
test_consumer_group test_topic  0           2              2              0      -          -          -

```

Figure 1.6: Kafka consumer group describe

1.6 Stop Broker and Check Cluster Health

```
1 #path-> /usr/odp/3.2.2.0-1/kafka/bin
2 ./kafka-topic.sh \
3 > --bootstrap-server stg-xxx:6667 \
4 > --describe \
5 > --topic test_topic
6
7 #zookeeper shell
8 ls /brokers/ids
```

```
root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdpsiddharth102:6667 --topic test_topic
Topic: test_topic    TopicId: 9pFEK8R2SD20cXHlhQ9dgg PartitionCount: 20      ReplicationFactor: 3      Configs: compression.type=producer,min.insync.replicas=1,segment.bytes=1073741824,max.message.bytes=1000000,index.interval.bytes=4096,retention.bytes=-1,segment.index.bytes=10485760
Partition: 0 Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
Partition: 1 Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
Partition: 2 Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
Partition: 3 Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
Partition: 4 Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1002
Partition: 5 Leader: 1002      Replicas: 1001,1002,1003      Isr: 1002,1003
Partition: 6 Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
Partition: 7 Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
Partition: 8 Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
Partition: 9 Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
Partition: 10 Leader: 1003     Replicas: 1003,1002,1001      Isr: 1003,1002
Partition: 11 Leader: 1003     Replicas: 1001,1003,1002      Isr: 1003,1002
Partition: 12 Leader: 1002     Replicas: 1002,1003,1001      Isr: 1002,1003
Partition: 13 Leader: 1003     Replicas: 1003,1001,1002      Isr: 1003,1002
Partition: 14 Leader: 1002     Replicas: 1001,1002,1003      Isr: 1002,1003
Partition: 15 Leader: 1002     Replicas: 1002,1001,1003      Isr: 1002,1003
Partition: 16 Leader: 1003     Replicas: 1003,1002,1001      Isr: 1003,1002
Partition: 17 Leader: 1003     Replicas: 1001,1003,1002      Isr: 1003,1002
Partition: 18 Leader: 1002     Replicas: 1002,1003,1001      Isr: 1002,1003
Partition: 19 Leader: 1003     Replicas: 1003,1001,1002      Isr: 1003,1002
```

Figure 1.7: Kafka describe after one broker down

```
WatchedEvent state:SyncConnected type:None path:null
[[zk: localhost:2181(CONNECTED) 0] ls /brokers/ids
[1002, 1003]
```

Figure 1.8: Broker id's in zookeeper

1.7 Again, produce and consume data for the topic test_topic

```
1 # Produce data
2 #path-> /usr/odp/3.2.2.0-1/kafka/bin
3 ./kafka-console-producer.sh \
4 > --bootstrap-server stg-XXX:6667 \
5 > --topic test_topic \
6 > --property parse.key=true \
7 > --property key.separator=":" \
8
9
10 # Consume data
11 #path-> /usr/odp/3.2.2.0-1/kafka/bin
12 ./kafka-console-consumer.sh \
13 > --bootstrap-server stg-xxx:6667 \
14 > --topic test_topic \
15 > --from-beginning \
16 > --group test_consumer_group \
17 > --property print.key=true \
18 > --property print.partition=true \
19 > --property print.offset=true \
20 > --property print.timestamp=true
```

```
[root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-console-producer.sh --bootstrap-server stg-hdpsiddharth102:6667 --topic test_topic --property parse.key=true --property key.separator=":"
>down1:xwcwc
>down2:cdcwc
>down3:wcwecw
>down6:cwvsv
>]
```

Figure 1.9: Kafka producer

```
[root@stg-hdpsiddharth104:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-console-consumer.sh --bootstrap-server stg-hdpsiddharth102:6667 --topic test_topic --from-beginning --group test_consumer_group --property print.key=true --property print.partition=true --property print.offset=true --property print.timestamp=true
CreateTime:1759229901853      Partition:6      Offset:0      down1      xwcwc
CreateTime:1759229905967      Partition:8      Offset:0      down2      cdcwc
CreateTime:1759229909544      Partition:6      Offset:1      down3      wcwecw
CreateTime:1759229921079      Partition:16     Offset:0      down6      cwvsv
```

Figure 1.10: Kafka consumer

1.8 Describe the topic again and observe the changes

```
1 #path-> /usr/odp/3.2.2.0-1/kafka/bin
2 ./kafka-topic.sh \
3 > --bootstrap-server stg-xxx:6667 \
4 > --describe \
5 > --topic test_
```

```

root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdpsiddharth102:6667
ic test_topic
Topic: test_topic      TopicId: 9pFEK8R2SD20cXH1hQ9dgg PartitionCount: 20      ReplicationFactor: 3      Configs: compression.type=producer
insync.replicas=1,segment.bytes=1073741824,max.message.bytes=1000000,index.interval.bytes=4096,retention.bytes=-1,segment.index.bytes=1048576
    Topic: test_topic      Partition: 0      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 1      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 2      Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 3      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 4      Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 5      Leader: 1002      Replicas: 1001,1002,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 6      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 7      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 8      Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 9      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 10     Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 11     Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 12     Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 13     Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 14     Leader: 1002      Replicas: 1001,1002,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 15     Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 16     Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 17     Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 18     Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 19     Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1002
root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]#

```

Figure 1.11: Kafka describe

1.9 Under-Replicated Partitions

```

1 #path-> /usr/odp/3.2.2.0-1/kafka/bin
2 ./kafka-topics.sh \
3 > --describe \
4 > --bootstrap-server stg-xxx:6667 \
5 > --under-replicated-partitions

```

Explanation

- Executes the `./kafka-topics.sh` script to `--describe` topics, connecting to the cluster via `--bootstrap-server` `stg-xxx:6667`. The command includes the `--under-replicated-partitions` flag, which acts as a filter to display only the topics that currently have partitions with fewer active replicas than their configured replication factor. This is a common command for diagnosing the health of a Kafka cluster.

```

root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdpsiddharth102:6667
ic test_topic --under-replicated-partitions
    Topic: test_topic      Partition: 0      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 1      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 2      Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 3      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 4      Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 5      Leader: 1002      Replicas: 1001,1002,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 6      Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 7      Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 8      Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 9      Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 10     Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 11     Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 12     Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 13     Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 14     Leader: 1002      Replicas: 1001,1002,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 15     Leader: 1002      Replicas: 1002,1001,1003      Isr: 1002,1003
    Topic: test_topic      Partition: 16     Leader: 1003      Replicas: 1003,1002,1001      Isr: 1003,1002
    Topic: test_topic      Partition: 17     Leader: 1003      Replicas: 1001,1003,1002      Isr: 1003,1002
    Topic: test_topic      Partition: 18     Leader: 1002      Replicas: 1002,1003,1001      Isr: 1002,1003
    Topic: test_topic      Partition: 19     Leader: 1003      Replicas: 1003,1001,1002      Isr: 1003,1002
root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]#

```

Figure 1.12: Kafka describe under-replica

1.10 Again, start the Kaka broker and observe the health

```

1 #path-> /usr/odp/3.2.2.0-1/kafka/bin
2 ./kafka-topic.sh \
3 > --bootstrap-server stg-xxx:6667 \
4 > --describe \
5 > --topic test_topic
6
7 #zookeeper shell
8 ls /brokers/ids

```

```

WatchedEvent state:SyncConnected type:None path:null
[[zk: localhost:2181(CONNECTED) 0] ls /brokers/ids
[1001, 1002, 1003]

```

Figure 1.13: Broker id's in zookeeper

```

root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdpsiddharth102:6667 --topic test_topic
Topic: test_topic    TopicId: 99FEK8R2SPD20cXHt0h9dgg PartitionCount: 20      ReplicationFactor: 3      Configs: compression.type=producer,min.insync.replicas=1,segment.bytes=1073741824,max.message.bytes=1000000,index.interval.bytes=4096,retention.bytes=-1,segment.index.bytes=10485760
Topic: test_topic    Partition: 0    Leader: 1002    Replicas: 1002,1003,1001    Isr: 1002,1003,1001
Topic: test_topic    Partition: 1    Leader: 1003    Replicas: 1002,1002,1003    Isr: 1002,1002,1003
Topic: test_topic    Partition: 2    Leader: 1001    Replicas: 1001,1003,1002    Isr: 1003,1002,1001
Topic: test_topic    Partition: 3    Leader: 1002    Replicas: 1002,1003,1001    Isr: 1002,1003,1001
Topic: test_topic    Partition: 4    Leader: 1003    Replicas: 1003,1001,1002    Isr: 1003,1002,1001
Topic: test_topic    Partition: 5    Leader: 1001    Replicas: 1001,1002,1003    Isr: 1002,1003,1001
Topic: test_topic    Partition: 6    Leader: 1002    Replicas: 1002,1001,1003    Isr: 1002,1003,1001
Topic: test_topic    Partition: 7    Leader: 1003    Replicas: 1003,1002,1001    Isr: 1003,1002,1001
Topic: test_topic    Partition: 8    Leader: 1001    Replicas: 1001,1003,1002    Isr: 1003,1002,1001
Topic: test_topic    Partition: 9    Leader: 1002    Replicas: 1002,1003,1001    Isr: 1002,1003,1001
Topic: test_topic    Partition: 10   Leader: 1003    Replicas: 1003,1002,1001    Isr: 1003,1002,1001
Topic: test_topic    Partition: 11   Leader: 1001    Replicas: 1001,1003,1002    Isr: 1003,1002,1001
Topic: test_topic    Partition: 12   Leader: 1002    Replicas: 1002,1003,1001    Isr: 1002,1003,1001
Topic: test_topic    Partition: 13   Leader: 1003    Replicas: 1003,1001,1002    Isr: 1003,1002,1001
Topic: test_topic    Partition: 14   Leader: 1001    Replicas: 1001,1002,1003    Isr: 1002,1003,1001
Topic: test_topic    Partition: 15   Leader: 1002    Replicas: 1002,1001,1003    Isr: 1002,1003,1001
Topic: test_topic    Partition: 16   Leader: 1003    Replicas: 1003,1002,1001    Isr: 1003,1002,1001
Topic: test_topic    Partition: 17   Leader: 1001    Replicas: 1001,1003,1002    Isr: 1003,1002,1001
Topic: test_topic    Partition: 18   Leader: 1002    Replicas: 1002,1003,1001    Isr: 1002,1003,1001
Topic: test_topic    Partition: 19   Leader: 1003    Replicas: 1003,1001,1002    Isr: 1003,1002,1001
Topic: test_topic    Partition: 20   Leader: 1001    Replicas: 1001,1002,1003    Isr: 1003,1002,1001
root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdpsiddharth102:6667 --topic test_topic --under-replicated-partitions

```

Figure 1.14: Kafka describe

1.11 Experimenting with test_topic_1 with a replication factor of 1

```

1 #path-> /usr/odp/3.2.2.0-1/kafka/bin
2 ./kafka-topics.sh \
3 > --create \
4 > --bootstrap-server stg-hdfsiddharth102:6667 \
5 > --replication-factor 1 \
6 > --partitions 3 \
7 --topic test_topic_1
8
9 ./kafka-topics.sh \
10 > --describe \
11 > --bootstrap-server stg-hdfsiddharth102:6667 \
12 > --topic test_topic_1
13
14 ./kafka-console-producer.sh \
15 >--bootstrap-server stg-hdfsiddharth102:6667 \
16 >--topic test_topia_1
17
18 kafka stop
19
20 ./kafka-console-consumer.sh \
21 > --bootstrap-server stg-hdfsiddharth102:6667 \
22 > --topic test_topic_1 \
23 > --from-beginning \
24 > --group test_consumer_group_1 \
25 > --property print.key=true \
26 > --property print.partition=true \
27 > --property print.offset=true \
28 > --property print.timestamp=true

```

```

root@stg-hdfsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh \
> --create \
> --bootstrap-server stg-hdfsiddharth102:6667 \
> --replication-factor 1 \
> --partitions 3 \
> --topic test_topic_1
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore '_' could collide. To avoid issues it is best to use either, but not both.
Created topic test_topic_1.
root@stg-hdfsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh \
> --describe \
> --bootstrap-server stg-hdfsiddharth102:6667 \
> --topic test_topic_1
Topic: test_topic_1    TopicId: RReLBWweQBardmlK2X6Bgw PartitionCount: 3      ReplicationFactor: 1    Configs: compression.type=producer,min.insync.replicas=1,segment.bytes=1073741824,max.message.bytes=1000000,index.interval.bytes=4096,retention.bytes=-1,segment.index.bytes=10485760
        Topic: test_topic_1    Partition: 0    Leader: 1002    Replicas: 1002  Isr: 1002
        Topic: test_topic_1    Partition: 1    Leader: 1001    Replicas: 1001  Isr: 1001
        Topic: test_topic_1    Partition: 2    Leader: 1003    Replicas: 1003  Isr: 1003
root@stg-hdfsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-console-producer.sh \
> --bootstrap-server stg-hdfsiddharth102:6667 \
> --topic test_topic_1
>hello
>siddharth on thw way!!
>hurray!!
>

```

Figure 1.15: New topic creation and producing data

```

root@stg-hdpsiddharth104:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-console-consumer.sh \
> --bootstrap-server stg-hdpsiddharth102:6667 \
> --topic test_topic_1 \
> --from-beginning
hello
siddharth on thw way!!
hurray!!

```

Figure 1.16: Consuming data from new topic

```

root@stg-hdpsiddharth102:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-topics.sh --describe --bootstrap-server stg-hdpsiddharth102:6667 --topic
ic test_topic_1
Topic: test_topic_1      TopicId: RReLBWweQBardmlK2X6Bgw PartitionCount: 3      ReplicationFactor: 1      Configs: compression.type=producer,min.
inSyncReplicas=1,segment.bytes=1073741824,max.message.bytes=1000000,index.interval.bytes=4096,retention.bytes=-1,segment.index.bytes=10485760
    Topic: test_topic_1      Partition: 0      Leader: 1002      Replicas: 1002      Isr: 1002
    Topic: test_topic_1      Partition: 1      Leader: none      Replicas: 1001      Isr: 1001
    Topic: test_topic_1      Partition: 2      Leader: 1003      Replicas: 1003      Isr: 1003

```

Figure 1.17: Kafka describe with one broker down

```

root@stg-hdpsiddharth104:/usr/odp/3.2.2.0-1/kafka/bin[nb6][stg]# ./kafka-console-consumer.sh --bootstrap-server stg-hdpsiddharth102:6667 --topic test_topic_1 --from-beginning --group test_c
onsumer_group_1 --property print.key=true --property print.partition=true --property print.offset=true --property print.timestamp=true
CreateTime:1759252592104      Partition:2      Offset:1      null      hello
CreateTime:1759252619539      Partition:2      Offset:1      null      hurray!!
CreateTime:1759253458984      Partition:2      Offset:2      null      bjkn
CreateTime:1759253458999      Partition:2      Offset:3      null      awesdf
CreateTime:1759253478744      Partition:0      Offset:4      null      bkkjjjjjjj
CreateTime:1759253480849      Partition:2      Offset:4      null      jijkik
CreateTime:1759253447129      Partition:0      Offset:9      null      huhi
CreateTime:1759253455694      Partition:0      Offset:1      null      nonono
CreateTime:1759253445884      Partition:0      Offset:2      null      vujib kn
CreateTime:1759253475047      Partition:0      Offset:3      null      bnknkn
CreateTime:1759253869167      Partition:0      Offset:4      user1      sid
CreateTime:1759253878254      Partition:2      Offset:6      user2      qwertty
CreateTime:1759253890438      Partition:2      Offset:7      user3      bioo

```

Figure 1.18: Kafka consumer with one broker down

Chapter 2

Task 2: Kafka Health Alert Setup

2.1 Bash Script for Alerts

```
1 #!/bin/bash
2 BROKERS_EXPECTED=3                                     # Total brokers in
3     your cluster
4 KAFKA_BIN_DIR=/usr/odp/3.2.2.0-1/kafka/bin          # Kafka binaries
5     path
6 ZOOKEEPER_HOST=stg-hdpsiddharth102:2181           # Zookeeper host
7 TOPICS=("test_topic" "test_topic_1")                  # List of topics to
8     check
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
```

```

24 echo "ALERT: $URP_COUNT under-replicated partitions found for
25   topic $topic" | tee -a /tmp/kafka_health.log
26 else
27 echo "No under-replicated partitions for topic $topic" | tee -a /
28   tmp/kafka_health.log
29 fi
30 done
31 }
32
33 check_brokers
34 check_urp

```

Explanation

'Variable Definitions (Lines 2-5):'

- **Lines 2-5:** These lines set up configuration variables for the script: the BROKERS_EXPECTED count (3), the path to Kafka's binaries (KAFKA_BIN_DIR), the Zookeeper connection string (ZOOKEEPER_HOST), and an array of TOPICS to be monitored.

'check_brokers' Function (Lines 8-18):

- **Line 8:** Defines a function to verify the number of active Kafka brokers.
- **Line 9:** Uses the Zookeeper client (`zkCli.sh`) to query Zookeeper and list the IDs of all currently registered brokers from the `/brokers/ids` znode.
- **Line 10:** Counts the number of active brokers found in the previous step.
- **Lines 11-17:** This block compares the actual number of brokers with the expected number. It prints and logs an ALERT to `/tmp/kafka_health.log` if the connection fails or if the broker count is too low; otherwise, it logs an OK status.

'check_urp' Function (Lines 20-29):

- **Line 20:** Defines a function to check for Under-Replicated Partitions (URP), which is a key sign of an unhealthy cluster.
- **Line 21:** Starts a loop that iterates through each topic listed in the TOPICS array.
- **Line 22:** For each topic, it runs `kafka-topics.sh` with the `--under-replicated-partitions` flag and counts the lines in the output. A count greater than zero indicates that the topic has URPs.
- **Lines 23-27:** If the URP count for a topic is greater than zero, it prints and logs an ALERT; otherwise, it logs a healthy status for that topic.

Execution (Lines 31-32):

- **Lines 31-32:** These lines call the two previously defined functions, `check_brokers` and `check_urp`, to execute the health checks.

2.2 Cron Job Setup

```

1 */1 * * * * /usr/local/bin/kafka_health_check.sh >> /var/log/
  kafka_health.log 2>&1

```

```
[root@stg-hdpsiddharth102:/usr/local/bin[nb6][stg]# tail -f /var/log/kafka_health.log
No under-replicated partitions for topic test_topic_1
ALERT: Some Kafka brokers are down! Expected 3, Found 2
ALERT: 20 under-replicated partitions found for topic test_topic
No under-replicated partitions for topic test_topic_1
ALERT: Some Kafka brokers are down! Expected 3, Found 2
ALERT: 20 under-replicated partitions found for topic test_topic
No under-replicated partitions for topic test_topic_1
ALERT: Some Kafka brokers are down! Expected 3, Found 2
ALERT: 20 under-replicated partitions found for topic test_topic
No under-replicated partitions for topic test_topic_1
Kafka brokers OK: 1001 1002 1003
No under-replicated partitions for topic test_topic
No under-replicated partitions for topic test_topic_1
```

Figure 2.1: Alert logs